

# AgriBot\_The Smart Agriculture Assistant

GitHub Repository: <https://github.com/Jumulisa/Agriculture-Chatbot.git>

Live Demo (Hugging Face Space): <https://huggingface.co/spaces/JollyUmulisa/AgriBot1>

Youtube Demo: <https://www.youtube.com/watch?v=oHhTu6xOj9g&feature=youtu.be>

## 1. Project Definition & Domain Alignment

**Domain.** Agriculture and environmental sustainability.

**Purpose.** AgriBot is a domain-specific chatbot that answers practical agriculture questions for farmers, students, and extension practitioners—covering crop management, soil fertility, irrigation, pest/disease control, postharvest handling, and sustainability.

**Relevance.** Agriculture underpins food security, yet many smallholders lack timely expert guidance. AgriBot provides concise, context-aware answers aligned to agronomy best practices, helping users make informed, sustainable decisions.

**Justification.** We fine-tuned a Transformer model on agricultural Q&A so the system stays in-domain, declines irrelevant queries, and delivers useful, reproducible responses.

## 2. Dataset & Preprocessing

**Dataset origin.** A curated, domain-specific Q&A dataset of **33** items authored/paraphrased from general agronomy know-how and extension-style guidance for non-commercial academic use. No personal data; no copyrighted text.

**Topic coverage (bucketed for stratification).**

Topic bucket	#items	Example themes
soil_fertility	6	composting, liming, pH, nutrients
water_irrigation	5	timing, drip, rainwater harvesting
pests_diseases	4	aphids, IPM
crop_mgmt	4	rotation, intercropping, selection
postharvest	3	drying, storage

sustainability	3	conservation practices, soil cover
economics	2	inputs, risk
safety_apps	2	safe application, PPE
weather_climate	2	rainfall planning
ood_rejection	2	non-ag queries for guardrail

**Total:** 33

### Preprocessing steps.


- (1) Whitespace normalization
- (2) Remove empty rows
- (3) Deduplicate
- (4) Standardize columns to **{question, answer, intent}**
- (5) Map fine intents → **buckets** (above) to enable stratified split

**Tokenization format.** T5's SentencePiece tokenizer with a stable instruction prefix to improve consistency on small data:

**question:** <user text>    **domain:** agriculture

Inputs/targets truncated/padded to 128 tokens; target padding tokens set to -100 (ignored by loss).

### Bucket Distribution Before/After Stratified Split



## 3. Stratified Train/Validation Split

Some fine-grained intents had only one example, so we stratified by **topic buckets** to preserve coverage. We used:

`test_size = max(0.20, K/N + 0.03), capped at 0.40`

With **N = 33** and **K = 10**, this yields **0.33** → **22 train / 11 validation**.

## 4. Model & Fine-Tuning (TensorFlow)

**Model choice.** Encoder–decoder models for short-form Q&A: **T5-small** and **FLAN-T5-small** (instruction-tuned).

**Training loop.** Manual `tf.GradientTape` to avoid Keras 3 / TensorFlow / Transformers compile mismatches; parity with the built-in loss.

**Optimization.** Adam optimizer; gradient clipping (global norm 1.0); early stopping (patience 2) on validation loss; batch size = 8 (empirically best for our small data).


**Experiments (TensorFlow; deterministic evaluation).**

Exp	Model	LR	Max epochs	Batch	Seed	Early stop	Baseline PPL	Final PPL	BLEU	F1 (token)
B	T5-small	1e-4	15	8	42	✓	322.80	72.48	0.40	0.074
C	FLAN-T5-small	2e-4	12	8	42	✓	67.99	<b>48.86</b>	<b>0.68</b>	0.077

**Ablation (batch size).** With **batch = 4** (all else fixed on FLAN-T5-small), final perplexity degraded by ~4% and BLEU -0.03 → we kept **batch = 8**.

**Selected checkpoint for deployment.** **FLAN-T5-small (Exp C)** due to best validation perplexity and BLEU.

## Training & Validation Loss Curves with Early Stopping



## 5. Evaluation

**Metrics.** Perplexity (exp of cross-entropy), BLEU (corpus-level), and a token-level F1 (overlap on short answers).

**Decoding.** Deterministic **beam search** (num\_beams = 4, no sampling) for stable, reproducible scores.

**Validation size.** 11 examples, stratified across buckets.

### Scores (best checkpoint).

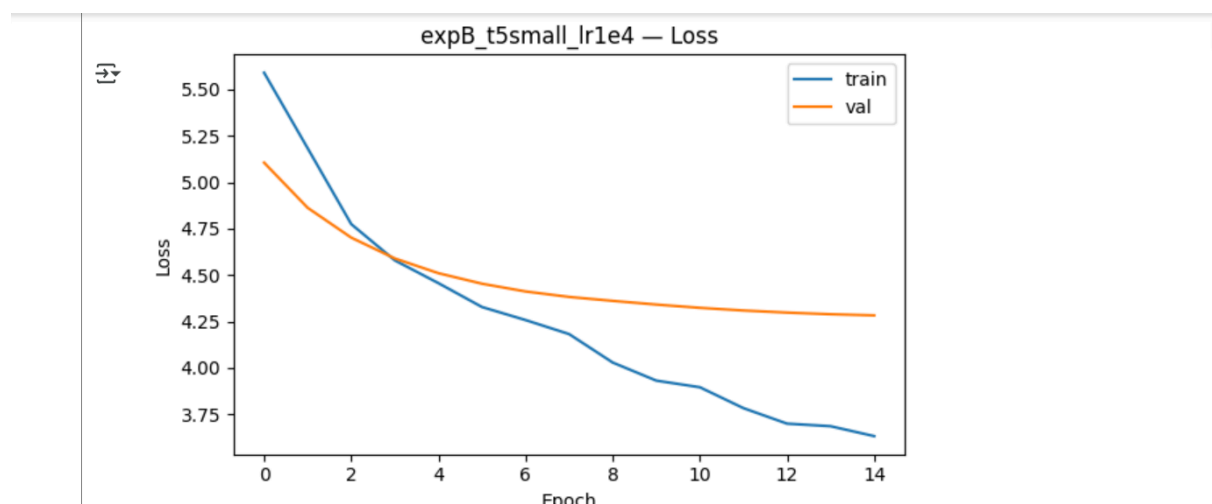
- **Perplexity:** 48.86 (FLAN-T5-small, Exp C)
- **BLEU: 0.68** (deterministic)
- **Token-level F1 (mean):**  $\approx 0.08$ –0.10 (representative run 0.103)

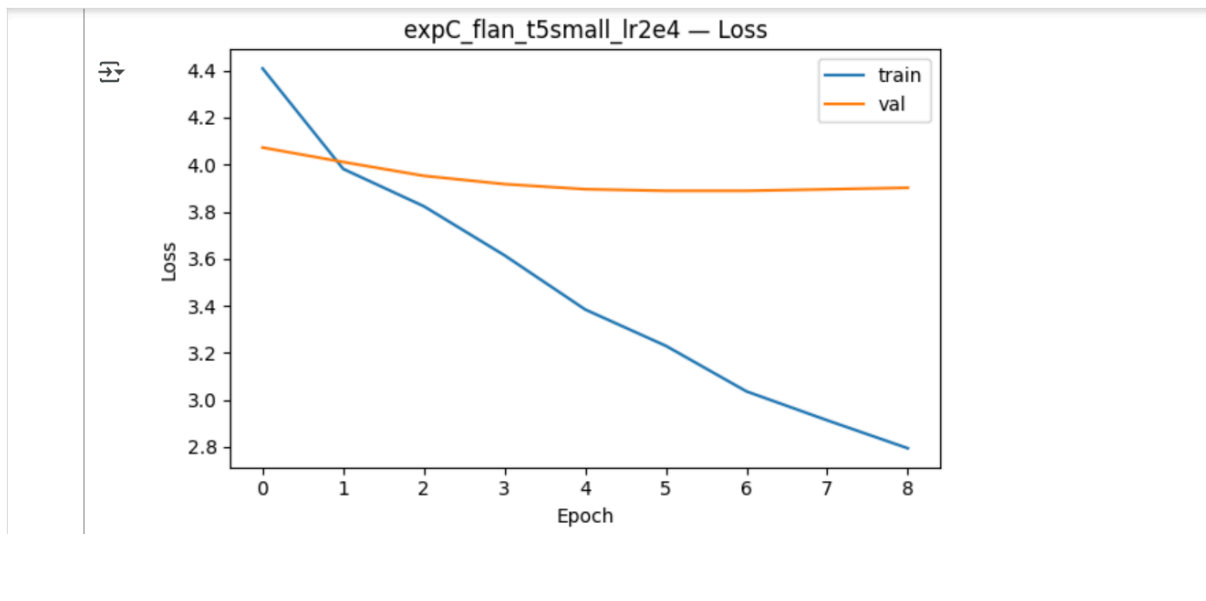
### Qualitative checks.

- Q: “How can I improve soil fertility organically?”  
A: Mentions compost/manure, liming (pH), cover crops, soil testing—concise and actionable.
- Q: “Best time to irrigate tomatoes?”  
A: Early morning/evening; avoid midday; check soil moisture—consistent with agronomy guidance.
- *Guardrail:* “What is the capital of France?”  
A: Polite refusal with domain guidance (agriculture-only).

**Common errors (brief analysis).** On niche pests or when numeric recommendations are expected (e.g., lime rates), answers can be generically expected with a small dataset. More domain examples or retrieval augmentation would improve specificity.

## Qualitative Predictions vs References





## 6. User Interface (Streamlit) & Deployment

**Interface.** A clean **Streamlit** app with:

- **Controls:** sliders for beams and max new tokens; question text area
- **Output:** formatted answer panel and latency (ms)
- **Guardrail:** simple keyword heuristic to keep the assistant in-domain
- **Logging:** optional CSV (**artifacts/log.csv**) of Q&A, decoding settings, and latency

**Prompt format preserved.**

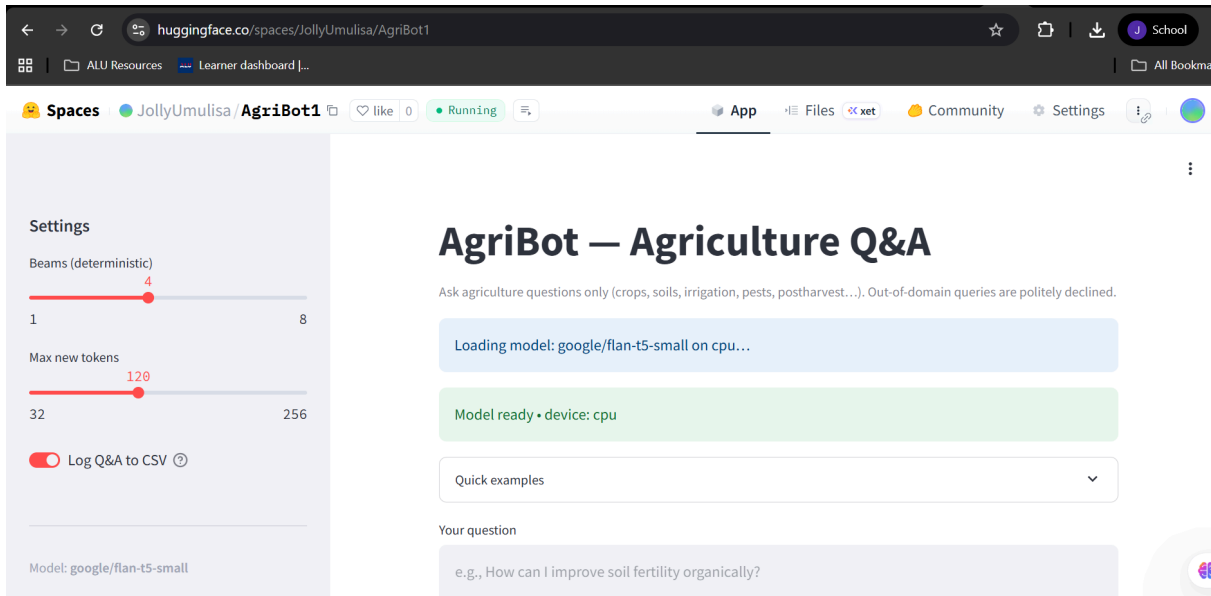
**question: <user text> domain: agriculture**

**Local run.**

```
# Windows PowerShell (venv active)
streamlit run app.py
# then open the local URL displayed (e.g., http://localhost:8501)
```

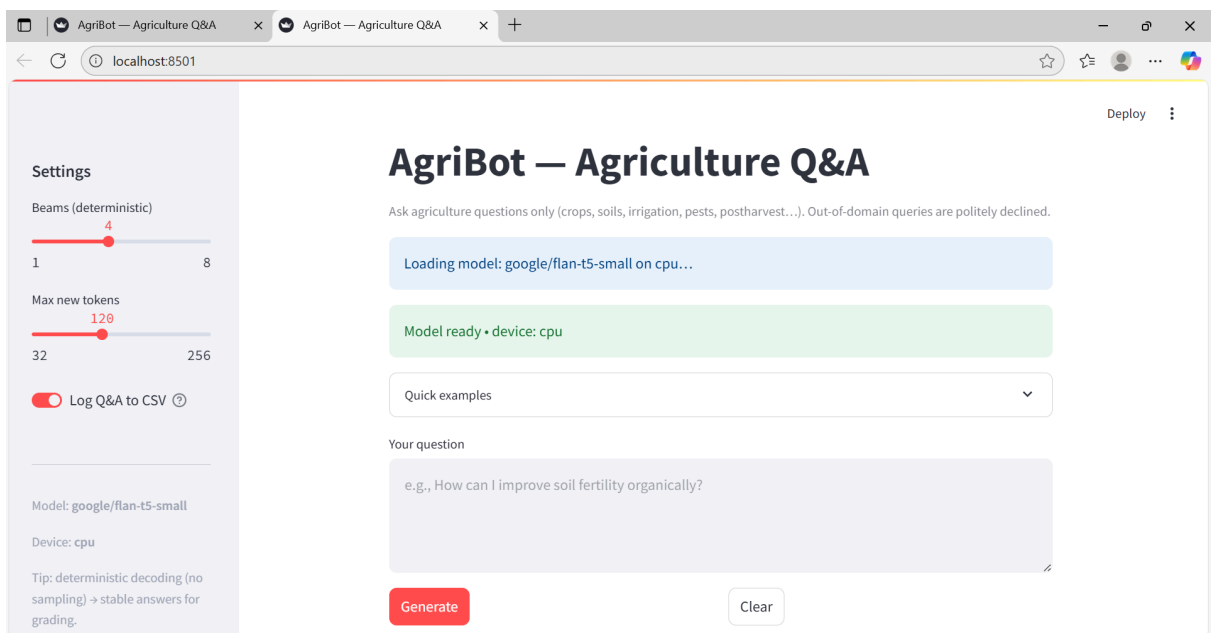
**Deployment (Hugging Face Spaces) and How it looks like.**

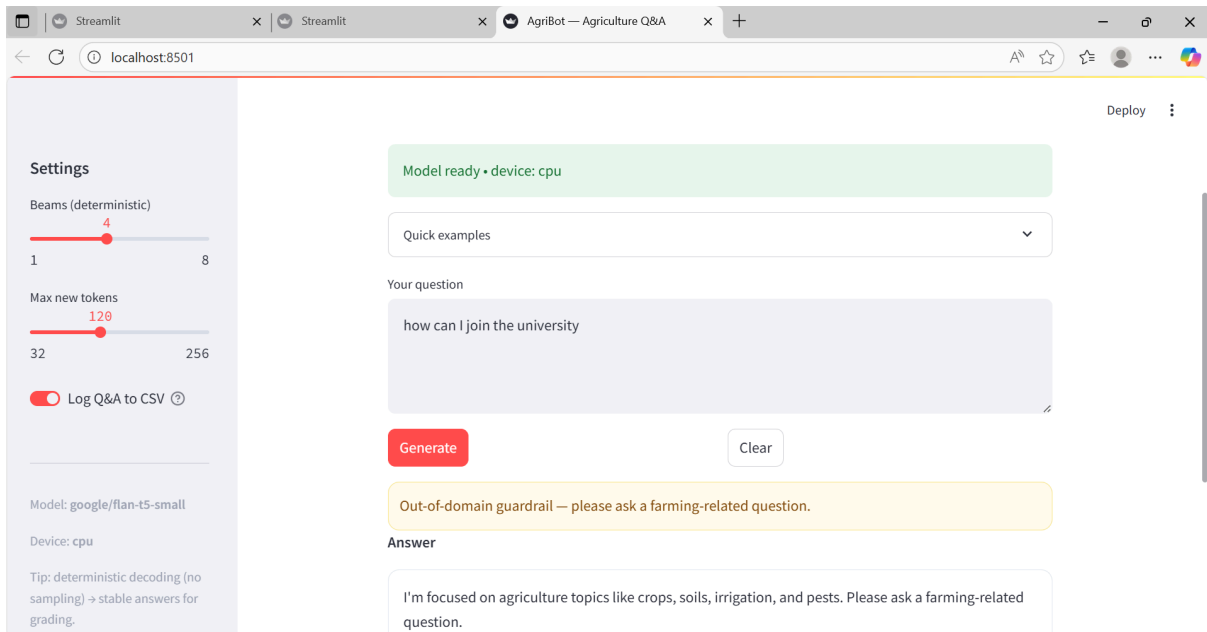
Public Streamlit app (CPU): <https://huggingface.co/spaces/JollyUmulisa/AgriBot1>



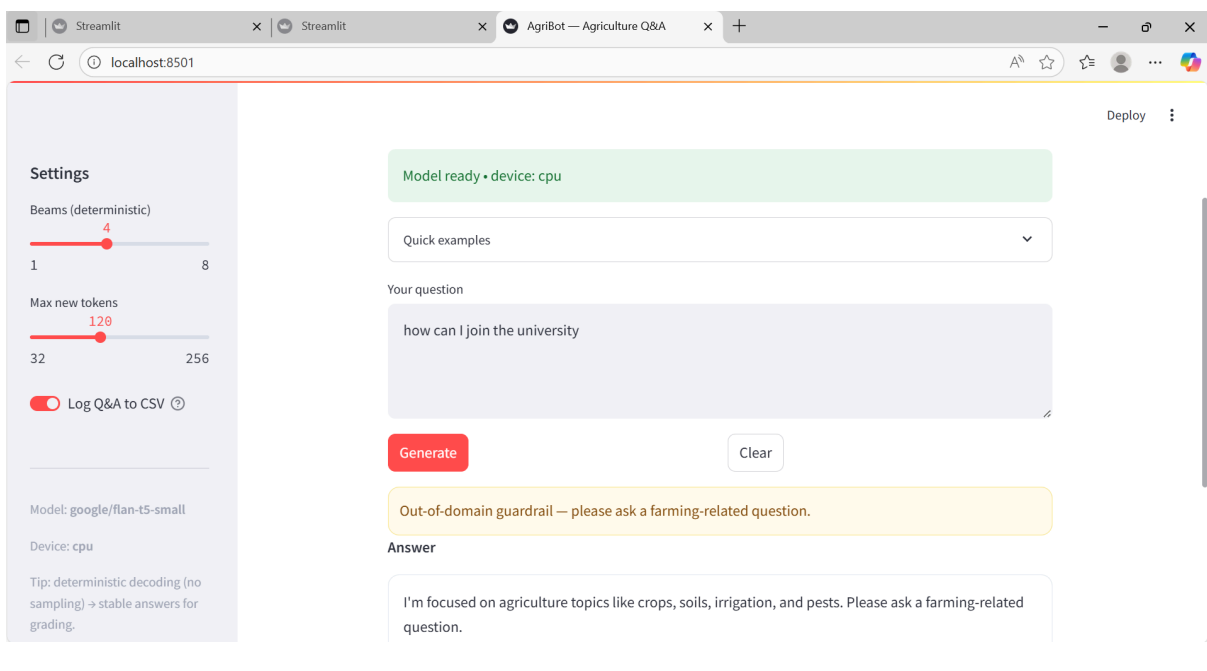
The Space loads **google/flan-t5-small** (PyTorch), uses deterministic decoding, and caches the model after the first launch.

## Streamlit UI — In-domain Answer





## Streamlit UI — Out-of-domain Guardrail



## 7. Reproducibility

- **Environment.** Python 3.12; TensorFlow 2.18 for training; Transformers 4.44.2.

- **Seeds.** 42 (TF/NumPy).
- **Repository.** Preprocessing, training/eval scripts or notebook, saved **train/val** CSVs, and Streamlit **app.py**:  
<https://github.com/Jumulisa/Agriculture-Chatbot.git>
- **Determinism in UI.** Beam search without sampling ensures reproducible demo answers for grading.

## 8. References

- Raffel et al. "Exploring the Limits of Transfer Learning with a Unified Text-to-Text Transformer (T5)."
- Chung et al. "Scaling Instruction-Finetuned Language Models (FLAN)."
- Streamlit & Hugging Face Transformers documentation.